



Installing OpenMPI

```
apt-get install libopenmpi-dev openmpi-bin  
openmpi-doc openmpi-common
```



Compiling and running MPI programs

`mpicc test.c`

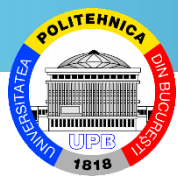
`mpirun -np 4 a.out`

`mpirun -np 3 date`

`./a.out`

Starts 4 processes.
Possibly on different machines.
They are identical but have different ids.
Works with non-MPI programs.

Works but starts only one process.



MPI example

```
#include<mpi.h>
```

```
#include<stdio.h>
```

```
int main(int argc, char * argv[])
```

```
{
```

```
    int rank;
```

```
    int nProcesses;
```

```
    MPI_Init(&argc, &argv);
```

```
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

```
    MPI_Comm_size(MPI_COMM_WORLD, &nProcesses);
```

```
    printf("Hello from %i/%i\n", rank, nProcesses);
```

```
    MPI_Finalize();
```

```
    return 0;
```

```
}
```



MPI example

```
#include<mpi.h>
```

```
#include<stdio.h>
```

```
int main(int argc, char * argv[])
```

```
{
```

```
    int rank;
```

```
    int nProcesses;
```

```
    MPI_Init(&argc, &argv);
```

```
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

```
    MPI_Comm_size(MPI_COMM_WORLD, &nProcesses);
```


```
    printf("Hello from %i/%i\n", rank, nProcesses);
```

```
    MPI_Finalize();
```

```
    return 0;
```

```
}
```

Start MPI Process





MPI example

```
#include<mpi.h>
```

```
#include<stdio.h>
```

```
int main(int argc, char * argv[])
```

```
{
```

```
    int rank;
```

```
    int nProcesses;
```

```
    MPI_Init(&argc, &argv);
```

```
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

```
    MPI_Comm_size(MPI_COMM_WORLD, &nProcesses);
```


```
    printf("Hello from %i/%i\n", rank, nProcesses);
```

```
    MPI_Finalize();
```

```
    return 0;
```

```
}
```

Get the id (rank)





MPI example

```
#include<mpi.h>
```

```
#include<stdio.h>
```

```
int main(int argc, char * argv[])
```

```
{
```

```
    int rank;
```

```
    int nProcesses;
```

```
    MPI_Init(&argc, &argv);
```

```
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

```
    MPI_Comm_size(MPI_COMM_WORLD, &nProcesses);
```


```
    printf("Hello from %i/%i\n", rank, nProcesses);
```

```
    MPI_Finalize();
```

```
    return 0;
```

```
}
```

**Get the total number of
processed**





MPI example

```
#include<mpi.h>
```

```
#include<stdio.h>
```

```
int main(int argc, char * argv[])
```

```
{
```

```
    int rank;
```

```
    int nProcesses;
```

```
    MPI_Init(&argc, &argv);
```

```
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

```
    MPI_Comm_size(MPI_COMM_WORLD, &nProcesses);
```

```
    printf("Hello from %i/%i\n", rank, nProcesses);
```

```
    MPI_Finalize();
```

```
    return 0;
```

```
}
```



**Print hello from all
processes.**



MPI example

```
#include<mpi.h>
```

```
#include<stdio.h>
```

```
int main(int argc, char * argv[])
```

```
{
```

```
    int rank;
```

```
    int nProcesses;
```

```
    MPI_Init(&argc, &argv);
```

```
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

```
    MPI_Comm_size(MPI_COMM_WORLD, &nProcesses);
```

```
    printf("Hello from %i/%i\n", rank, nProcesses);
```

```
    MPI_Finalize();
```

```
    return 0;
```

```
}
```

**Stop the MPI
environment.**



MPI example executed

```
#include<mpi.h>
#include<stdio.h>
```

```
int main(int argc, char * argv[])
{
    int rank;
    int nProcesses;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &nProcesses);
    printf("Hello from %i/%i\n", rank, nProcesses);
    MPI_Finalize();
    return 0;
}
```

Hello from 0/4

```
#include<mpi.h>
#include<stdio.h>
```

```
int main(int argc, char * argv[])
{
    int rank;
    int nProcesses;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &nProcesses);
    printf("Hello from %i/%i\n", rank, nProcesses);
    MPI_Finalize();
    return 0;
}
```

Hello from 3/4

```
#include<mpi.h>
#include<stdio.h>
```

```
int main(int argc, char * argv[])
{
    int rank;
    int nProcesses;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &nProcesses);
    printf("Hello from %i/%i\n", rank, nProcesses);
    MPI_Finalize();
    return 0;
}
```

Hello from 2/4

```
#include<mpi.h>
#include<stdio.h>
```

```
int main(int argc, char * argv[])
{
    int rank;
    int nProcesses;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &nProcesses);
    printf("Hello from %i/%i\n", rank, nProcesses);
    MPI_Finalize();
    return 0;
}
```

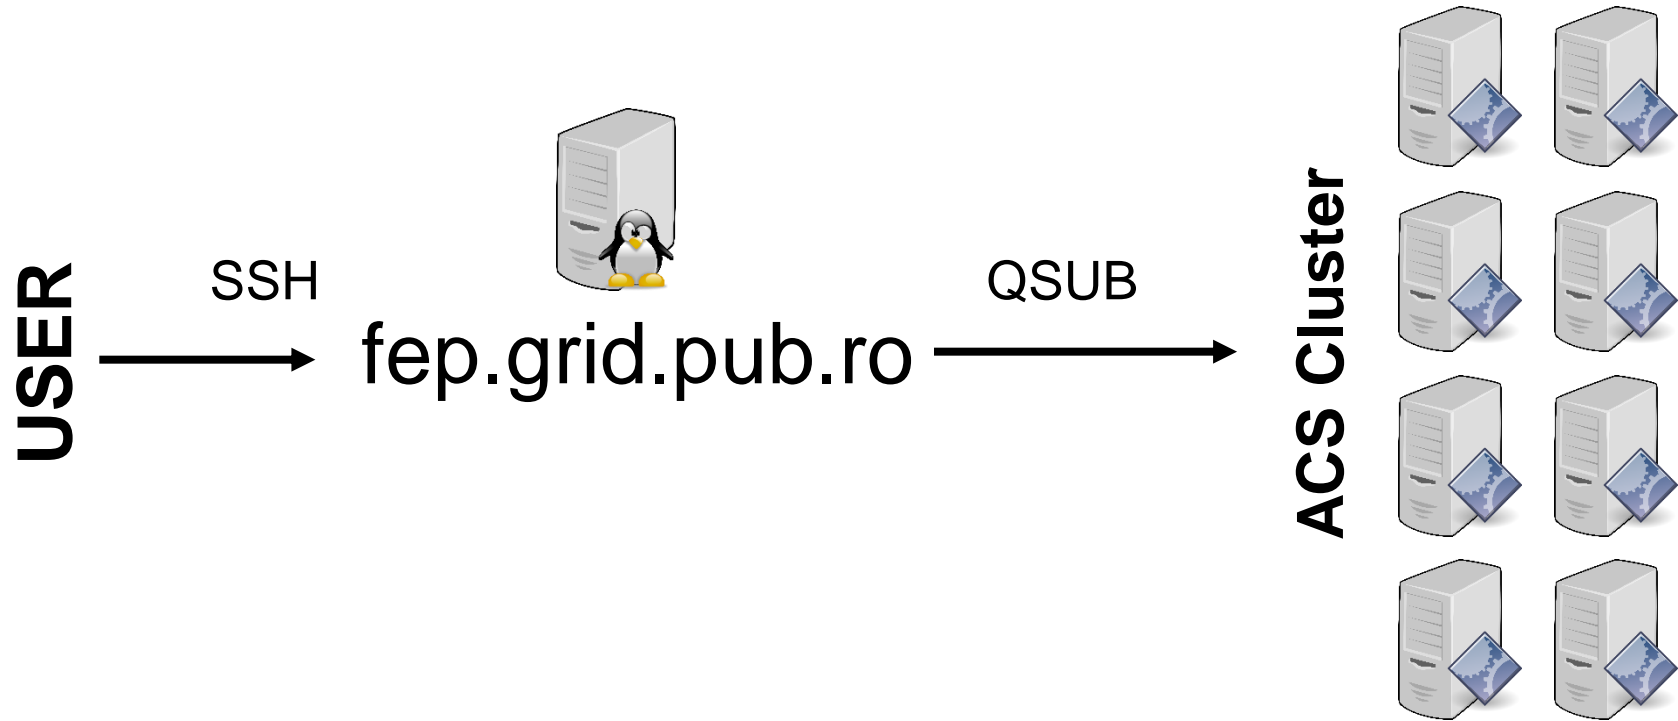
Hello from 1/4



MPI memory

- There is no such thing as shared memory in MPI
- All variables are local per process (equivalent to private in OpenMP)
- To get information from one process to the other you have to use explicit communication
 - Send/Recv
 - Broadcast
 - Scatter
 - Gather

Running on the cluster





Account

cs.curs.pub.ro
Account

name.familyName



QSUB

Parameters

- cwd (change working directory)**
- q queue_name (select a queue)**



Run on fep.grid.pub.ro:

```
qsub -cwd -pe openmpi*1 <nTasks> <<EOF  
> module load libraries/openmpi-1.6-gcc-4.7.0  
> mpirun -n <nTasks> ./executable_name  
> EOF
```

Current Working Directory

Makes sure the input/output files are taken/placed in the current directory.



MPI@ACS

Run on fep.grid.pub.ro:

```
qsub -cwd -pe openmpi*1 <nTasks> <<EOF  
> module load libraries/openmpi-1.6-gcc-4.7.0  
> mpirun -n <nTasks> ./executable_name  
> EOF
```

Specify how many machines/tasks to use



MPI@ACS

Run on fep.grid.pub.ro:

```
qsub -cwd -pe openmpi*1 <nTasks> <<EOF  
> module load libraries/openmpi-1.6-gcc-4.7.0  
> mpirun -n <nTasks> ./executable_name  
> EOF
```

Bash shorthand tells qsub a script is following
The script is surrounded by EOF



Using a bash script bash script:

```
qsub -cwd -pe openmpi*1 <nTasks> script.sh
```

script.sh contents:

```
#!/bin/bash
```

```
module load libraries/openmpi-1.6-gcc-4.7.0
```

```
mpirun -n <nTasks> ./executable_name
```